

Analysis: Theme Park

At first glance, this problem appears to be a straightforward simulation: keep adding groups until you run out space on the roller coaster (or run out of groups), then re-form the queue and start again. Repeat this **R** times and you're done.

For the Small, that was enough; and we got more than a few questions during the contest from contestants thinking that they were unable to solve the Large because their computers were so slow. The fact is that, with limits so large -- up to 10^3 groups queuing for 10^9 rides -- you need to come up with a smarter algorithm to solve the problem, since that one is $O(NR)$.

Optimization One

When you're sending the roller coaster out for 10^9 rides, you've got to expect that a few of them are going to be the same. If you store the queue of groups as an unchanging array, with a pointer to the front, then every time you send out a ride *s*, you could make a note of where it ended, given where it started. Then the next time you see a ride starting with the same group in the queue, you can do a quick lookup rather than iterating through all the groups.

That speeds up the algorithm by a factor of 10^3 in the worst case, leaving us with $O(R)$ operations. There are some other ways of speeding up the calculation for any given roller coaster run: for example, you could make an array that makes it $O(1)$ to calculate how many people are in the range [group_a, group_b] and then binary search to figure out how many groups get to go in $O(\log(N))$ time. That gives a total of $O(R \log N)$ operations.

Optimization Two

As we observed in Optimization One, you're going to see a lot of repetition between rides. You're also going to see a lot of repetition between groups of rides. In the example in the problem statement, the queue was made up of groups of size [1, 4, 2, 1]. 6 people get to go at once. Let's look at how the queue changes between rides:

```
1, 4, 2, 1 [5]
2, 1, 1, 4 [4]
4, 2, 1, 1 [6]
1, 1, 4, 2 [6]
2, 1, 1, 4 [4]
4, 2, 1, 1 [6]
1, 1, 4, 2 [6]
```

As you may have noticed, there's a *cycle* of length three: starting from the second run, every third queue looks the same. We make 16 Euros when that happens, which means we'll be making 16 Euros every 3 runs until the roller coaster stops rolling.

So if the roller coaster is set to go 10^9 times: the first time it makes 5 Euros; then there are 999999999 runs left; and every three of those makes 16 Euros. 3 divides 999999999 evenly -- if it didn't, we'd have to do some extra work at the end -- so we make $5 + (999999999 / 3 * 16) = 5333333333$ Euros in total.

It turns out that a cycle *must* show up within the first $N+1$ rides, because there are only **N** different states the queue can be in (after **N**, you have to start repeating). So you only have to

simulate **N** rides, each of which takes $O(N)$ time in the worst case, before finding the cycle: that's an $O(N^2)$ solution.

Optimization Three

Either of the optimizations above should be enough. But if you're a real speed demon, you can squeeze out a little more efficiency by combining the binary search that we mentioned briefly in Optimization One with the cycle detection from Optimization Two, bringing our running time down to $O(N \log N)$. An alternate optimization can bring us down to $O(N)$; we'll leave that as an exercise for the reader. Visit our [Google Group](#) to discuss it with the other contestants!