

Analysis: Security Update

Let R_i be the number of computers that receive the update before computer i , and T_i be the time between computer 1 and computer i receiving the update. For each i , the input gives us exactly one of these numbers. We can set $R_1 = T_1 = 0$ for convenience.

A simplified problem

Let us assume for now that we have all the T_i s. If computers i and j share a direct connection and $T_i = T_j$, then any path that gets to computer i in time T_i does not go through computer j , and vice versa, because all latencies are positive. Therefore, we can assign any positive latency to all those connections. If computer i has a given T_i , it means that any connection coming from computer j with $T_j < T_i$ needs to have a latency of at least $T_i - T_j$, or otherwise the update could get to computer i in less than T_i time by getting to computer j in T_j time and then using that connection. In addition, for at least one j , the latency of the connection between i and j has to be exactly $T_i - T_j$, or otherwise the time for the update to reach computer i would be larger than T_i . One simple way to solve this problem is to make all connections between computers with different T values have a latency of exactly $|T_i - T_j|$; this takes $O(\mathbf{D})$ time.

Notice that the algorithm above finds a valid assignment for any set of T_i s. To solve the actual test sets, we are left with the problem: given some T_i s and some other R_i s, assign all of the non-given values in a way such that sorting the computers by T_i leaves them sorted by R_i , and vice versa. In particular, computers with equal T values should have equal R values, and vice versa.

Test Set 1

In this test set, we can solve the subproblem from the previous section by setting $T_i := R_i$.

Test Set 2

For Test Set 2, we again focus on solving the subproblem. We do that by first ordering the computers by what is going to be their final T_i value (or equivalently, by their final R_i value). We can partition the set of computers other than the source computer into two: those for which we know R_i (part R) and those for which we know T_i (part T). We can sort each of those in non-decreasing order of the known value. We now have 2 sets that are in the right relative order, and need to merge them as in the last step of [Merge Sort](#). We assign the source computer first. Then we iterate through the remaining $\mathbf{C}-1$ slots in order. Suppose we have already merged N computers, and let computer k be the last one of those. Let i and j be the first computers remaining in parts R and T, respectively. If $R_i \leq N$, we take computer i next and assign $T_i := T_k$ if $R_i = R_k$, and $T_i := T_{k+1}$ otherwise. If $R_i > N$, we take computer j next and assign $R_j := R_k$ if $T_j = T_k$ and $R_j := N$ otherwise.

We can prove that if the original set of values is consistent with at least one latency assignment (which the statement guarantees), this generates a valid order and assignment of missing values, and moreover, it generates one in which the T value of the last computer in the order is minimal. We do that by induction on the number of computers. For a single computer, this is trivially true. Suppose we have $\mathbf{C} > 1$ computers. By our inductive hypothesis, the first $\mathbf{C}-1$ computers in the order were ordered and assigned values in a consistent way, with a minimal T value for the last computer among all options. Let us say the last computer in the full order is

computer i , and the next-to-last computer is computer j . By definition of how we assign missing values, $R_i = R_j$ if and only if $T_i = T_j$. If indeed $R_i = R_j$ and $T_i = T_j$, then the condition for the final assignment is equivalent to the inductive hypothesis. If computers i and j come from the same part, then the ordering choice between them was fixed, and the assignment of T values if needed is clearly minimal. So consider further the case in which computer i comes from a different part than computer j , and their R and T values are different. We have two cases: either computer i was in part R , or in part T .

If computer i was in part R , then its assigned T value is by definition the largest among all computers, and it's the smallest possible for it to go after computer j , whose value is minimal by the inductive hypothesis. As for the order, $R_i \leq C-1$ per the limits. Since computer j comes from part T and was chosen for position $C-1$ (when N was $C-2$), that means $R_j > C-2$. Therefore, $R_i = C-1$, and the chosen position is correct.

If, on the other hand, computer i was in part T , then its T value is minimal because T_i is fixed. As for the order, notice that all computers have either a T value strictly less than T_i or an R value strictly less than $C-1$, so none of them could have been last. By the inductive hypothesis, T_j is minimal among all possible orders, which means, by the existence of a full assignment, it has to be $T_j < T_i$, which implies the consistency of the final order and value assignment.