

Analysis: Palindromic Sequence

To solve the Small dataset, we will build up the palindrome one letter at a time. Suppose we know the first X letters of the K th lexicographically least palindrome. Call this prefix of the answer S . We can guess the next letter and see how many palindromes there are with the new prefix (we will address this computation shortly). If there are more than K , then we know we need to guess a lexicographically smaller letter. Let $P(S)$ denote the number of palindromes of Hannah's language that have S as a prefix. Then we are looking for the lexicographically largest character c such that $P(S + c) \leq K$. Naively, we have to make $O(NL)$ guesses.

So how do we actually calculate the number of possible continuations? Fortunately, the bounds are Small enough that we don't need to do anything sophisticated. For every possible length M , check and see if the current prefix S is consistent with a palindrome of length M . If it is, then there are L^D possible palindromes of length M , where $D = \max(0, \text{floor}((M+1)/2) - X)$. d represents the number of characters that can be freely chosen (the rest are fixed to keep the string a palindrome). This can be done naively in $O(N^2)$ time. In total, the Small can be solved in $O(N^3L)$ time.

To solve the Large dataset, notice that when $K \leq N$, the answer is a string that consists of the letter a repeated K times. When $K > N$, it's important to notice that the answer is very close to N .

For example, consider the case when $N = 10^5$ and $L = 2$. How many palindromes begin and end with 49900 letter as ? To make a rough count, let's consider only the palindromes with length exactly N . Of the 200 letters between the as , the first 100 are unconstrained, while the last 100 have to match the corresponding letter in the first 100. In total this is 2^{100} different palindromes, which far exceeds that maximum bound for K (10^{18} is about 2^{60}).

The example above can be generalised to state that whenever $K > N$, the K th palindrome will begin and end with at least $N - \log_L(K)$ copies of the letter a , leaving at most $2 \cdot \log_L(K)$ letters in the middle that might not be the letter a . The algorithm in the Small dataset can be adapted to solve the Large dataset after these observations have been made. $2 \cdot \log_L(K)$ is no more than 100, so this has a similar running time to the Small dataset.